

# AIND Project 2: Building a Game-Playing Agent

by Dmitri Lihhatsov / GitHub: @dlihhats

## Introduction

The goal of this project is to create an AI agent for a game of Isolation. In our version of the game, players can only move to cells that lie two by one cells away from the player's current location – just like a knight figure in the game of chess. The last player who cannot move any further loses the game.

Our AI agent tries to find an optimal move by traversing down the tree of all possible moves using a Minimax search algorithm with Alpha-Beta pruning and iterative deepening given a time constraint for every move. Our goal is to come up with an heuristic function describing a state of the game that would outperform another predefined AI agent. This predefined AI agent uses the same search method but with a heuristic function equal to the difference between the number of possible player moves and opponent moves.

## Evaluation Method

We will evaluate the performance of three heuristic functions  $f(\text{game})$  in comparison with the predefined AI agent by making them all play 100 games against the following AI opponents:

1. **AB\_Improved**: Minimax, Alpha-Beta, iterative deepening with  $f(\text{game}) = \text{my\_moves} - \text{opponents\_moves}$
2. **AB\_Center**: Minimax, Alpha-Beta, iterative deepening with heuristic function evaluating the distance to the centre of the board.
3. **AB\_Open**: Minimax, Alpha-Beta, iterative deepening with  $f(\text{game}) = \text{my\_moves}$
4. **MM\_Improved**: Minimax with search depth equal to 3 and  $f(\text{game}) = \text{my\_moves} - \text{opponents\_moves}$
5. **MM\_Center**: Minimax with search depth equal to 3 and heuristic function evaluating the distance to the centre of the board.
6. **MM\_Open**: Minimax with search depth equal to 3 and  $f(\text{game}) = \text{my\_moves}$
7. **Random**: An agent that selects its next move randomly.

## Heuristic Functions

In order to come up with an heuristic function that would be superior than a simple difference between number of player's moves and opponent's moves, we have to:

- keep the heuristic as simple as possible to allow the algorithm search deeper in the allotted time
- improve the pruning of the tree by highlighting the difference between various game states and finding equivalent game states.

Thus, we evaluated the following heuristic functions.

1. **AB\_Custom**: This agent uses a function that capitalises on the successful difference between player's moves and opponent's moves but also considers the ratio. We reckoned it to be more accurate than simply calculating the difference because it distinguishes various game states that the AB\_Improved agent finds equivalent. For example, considering two game states, 1) when the player has 5 moves and the opponent has 4 moves; and 2) when the player has 8 moves and the opponent has 7 moves; AB\_Custom agent will prefer 1) while AB\_Improved considers them to be equivalent.

$$f = \text{own\_moves} - \text{opp\_moves} + \text{own\_moves} / \text{opp\_moves}$$

2. **AB\_Custom\_2**: This agent uses a function that evaluates the ratio between the number of player's moves and opponent's moves. We wanted to compare its performance with a slightly more complex heuristic function of AB\_Custom.

$$f(\text{game}) = \text{own\_moves} / \text{opp\_moves}$$

3. **AB\_Custom\_3**: With this agent we decided to try out a function that would concentrate at the top left part quadrant of the board. We wanted to check an assumption that focusing on a particular part of the board could bring an advantage. We didn't have much hope for this function but as we will see, it still performed quite well.

$$f = \text{own\_moves\_in\_top\_left} - \text{opp\_moves\_in\_top\_left}$$

## Evaluating Performance

We run 100 games between all opponents with each opponent making the first move in 50 games and randomly determining the initial location of the players. These games are independent, and we will run a Z-test to determine if the difference between win rates of different agents is statistically significant.

	$X_1$	$X_2$
1	a	b
0	c	d
$\Sigma$	$n_1$	$n_2$

$$\hat{p}_1 = \frac{a}{n_1}$$

$$\hat{p}_2 = \frac{b}{n_2}$$

$$\text{Confidence interval for } p_1 - p_2: \hat{p}_1 - \hat{p}_2 \pm z_{1-\frac{\alpha}{2}} \sqrt{\frac{\hat{p}_1(1-\hat{p}_1)}{n_1} + \frac{\hat{p}_2(1-\hat{p}_2)}{n_2}}$$

$$Z\text{-test} : Z(X_1, X_2) = \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{P(1-P)(\frac{1}{n_1} + \frac{1}{n_2})}}$$

$$P = \frac{\hat{p}_1 n_1 + \hat{p}_2 n_2}{n_1 + n_2}$$

## Results

\*\*\*\*\*

### Playing Matches

\*\*\*\*\*

Match #	Opponent	AB_Improved		AB_Custom		AB_Custom_2		AB_Custom_3	
		Won	Lost	Won	Lost	Won	Lost	Won	Lost
1	AB_Improved	48	52	55	45	56	44	52	48
2	AB_Center	57	43	55	45	58	42	52	48
3	AB_Open	47	53	54	46	54	46	48	52
4	MM_Improved	87	13	81	19	89	11	80	20
5	MM_Center	91	9	95	5	93	7	92	8
6	MM_Open	88	12	88	12	92	8	88	12
7	Random	98	2	97	3	99	1	100	0
-----									
Win Rate:		73.7%		75.0%		77.3%		73.1%	

### Z-test

#### AB\_Improved vs. AB\_Custom

95% confidence interval for a win probability, AB\_Improved: [(0.703, 0.768)]

95% confidence interval for a win probability, AB\_Custom: [(0.717, 0.781)]

95% confidence interval for the difference: [(-0.059, 0.033)]

p-value, one-sided, AB\_Custom is better than AB\_Improved: 0.291

p-value, two-sided, AB\_Custom is different from AB\_Improved: 0.582

#### AB\_Improved vs. AB\_Custom\_2

95% confidence interval for a win probability, AB\_Improved: [(0.703, 0.768)]

95% confidence interval for a win probability, AB\_Custom\_2: [(0.740, 0.802)]

95% confidence interval for the difference: [(-0.081, 0.009)]

**p-value, one-sided, AB\_Custom\_2 is better than AB\_Improved: 0.060**

p-value, two-sided, AB\_Custom\_2 is different from AB\_Improved: 0.120

## Conclusion

Z-test shows that unfortunately with the level of confidence at 95% none of our heuristic functions are statistically any better than the AB\_Improved agent. However, with a confidence level at 90% we see that the AB\_Custom\_2 agent with an heuristic function  $f(\text{game}) = \text{own\_moves} / \text{opp\_moves}$  performs better than AB\_Improved (p-value = 0.060).

Therefore, for an intelligent agent in our game of Isolation, we recommend using a minimax algorithm with alpha-beta pruning, iterative deepening and an heuristic function equal to the ratio between player's moves and opponent's moves for the following reasons:

1. This heuristic function is simple enough to calculate under time restrictions to reach significant depths in the search tree.
2. This heuristic function amplifies the difference between various game states better than a function that calculates the difference between the number of moves. Thus, it allows to select the next move more intelligently.
3. Variability in the scores of different game states provides efficient alpha-beta pruning of the search tree allowing even deeper searches.